

Proofs for Now and Proofs for Later

The Amortization Horizon of Human and Artificial Reasoning in Lean

Working paper

2 August 2026

Abstract

A checked proof establishes that a proposition follows, but not what future it was built to serve. We propose the *amortization horizon*: the range over which a constructor expects an abstraction to repay its cost. A proof consequently has three separable products—a kernel certificate, working memory, and an interface for later proofs and readers. We test this distinction in 3,635 pairs of validated Lean proofs of identical statements. AI-provenance scripts create more named local claims, but each claim receives fewer explicit downstream references, crosses fewer later claim boundaries, and is far more often generically named. A scope-correct analysis of elaborated terms confirms more dead AI binders, yet finds equal rates of multi-use binders: tactics can turn disposable source state into kernel sharing. An exploratory Lean-model analysis finds a short information pulse at human claim boundaries, strongly attenuated at AI boundaries. These results locate the divergence not in deductive power or biological kind, but in which consumer the construction serves. Recent machine refactoring and library-building systems support the corresponding prediction: changing the horizon changes the artifact. We formalize the idea as amortized proof complexity over theorem streams, a quantity that single-theorem benchmarks cannot observe.

1 The missing future of a checked proof

A theorem prover settles a sharply posed question: did this term check at this type? That question is essential, but it is not the whole function of proof. A mathematician also uses a proof to find the invariant that makes a result intelligible, to remember a route through an argument, to expose a lemma that will solve the next theorem, and to leave an object another person can repair (Rav, 1999). These uses extend beyond the instant of verification. Two certificates of the same statement may thus agree perfectly as kernel inputs while serving different futures.

This paper argues that this future is the most useful level at which to compare present human- and AI-provenance Lean proofs. The common contrast—human proof versus machine proof—encourages an essentialist question: what signature does the identity of the constructor leave in a proof graph? That question is poorly controlled. Theorem, library, tactic, model, and representation all change graph statistics. More importantly, it mistakes an objective for an organism. Most generated proofs are rewarded when the current declaration compiles. Much human formalization occurs inside a longer activity of exposition, maintenance, and library construction; early workflow evidence also finds that users want AI assistance while preserving high-level control (Collins et al., 2026). Different artifacts are therefore rational even before any difference in cognitive machinery is posited.

We call the relevant variable the *amortization horizon*: how far into future reasoning the constructor expects a present abstraction to pay. At a short horizon, a local claim can be useful merely because it records a search result for the next tactic. Its name need only be distinct; its statement need not be reusable; it may disappear under elaboration. At a long horizon, names,

boundaries, generality, and stable dependencies become investments. Their up-front cost can be recovered across later steps, theorems, versions, and readers.

This account makes four contributions. First, it separates three products of formal construction: certificate, episode structure, and interface. Second, it tests their separation in 3,635 exact-statement human–AI pairs, with a scope-aware kernel-term validation on 298 pairs. Third, it offers an exploratory information-theoretic test inspired by work on how speakers regulate surprise in conversation (Bergey and DeDeo, 2024). Finally, it proposes a sequence-level object, amortized proof complexity, that charges for a library and the stream of proofs it enables. This reframes a central evaluation problem: success at one theorem measures certificate production, not mathematical accumulation.

The argument deliberately narrows several earlier claims in this repository. Viteri and DeDeo’s proof-network work establishes that proof organization can support robust belief and reveals heavy-tailed reuse (Viteri and DeDeo, 2022). Our replication found those results informative but highly sensitive to the graph boundary. Near-saturated belief did not distinguish paired certificates, and a newly discovered de Bruijn-scope conflation makes the inherited unexpanded-term topology unsuitable as a semantic reuse measure. Rather than force authorship into those statistics, we use the failure to identify the missing object: who or what must use the proof next?

2 A proof has three consumers

The word “proof” compresses at least three artifacts.

Certificate. The kernel consumes an elaborated term and checks that it inhabits the target type. The certificate may contain local `let` binders and shared subexpressions, but the kernel does not require human-readable names, explanatory order, or a useful public API.

Episode structure. The constructor and an immediate reader consume the source script during the current attempt. A named `have` makes a fact addressable, reduces the active goal, and records search state. It can perform this role even if no later source token names it: a contextual tactic such as `linarith` may consume it implicitly, or elaboration may erase it.

Interface. Future proofs and maintainers consume exported declarations, stable names, general statements, and explanatory boundaries. An interface is not just a set of dependencies. It is a retrieval scheme: it partitions a derivation into units that can be recognized and invoked without reconstructing their interiors. DeDeo and Duede call attention to the associated correspondence problem: a formal derivation can exist while failing to track the proof idea that is supposed to justify its relevance (DeDeo and Duede, 2026).

These consumers define different notions of reuse. A proof-term binder used twice is a sharing node in a directed acyclic graph. A source claim referenced after several later claims is working memory with temporal reach. An exported theorem cited by later files is persistent memory. None is a universal measure of modularity. In particular, raw occurrence counts can mistake tactic expansion for conceptual importance, while a theorem that is used only once may still identify the right explanatory cut.

For an abstraction z , a minimal decision model is

$$V_\gamma(z) = \sum_{u \in U(z)} \gamma^{d(z,u)} s(u) - c_{\text{make}}(z) - c_{\text{interface}}(z), \quad (1)$$

where $U(z)$ is the set of anticipated uses, $s(u)$ is work saved at use u , and d is distance in steps, theorems, time, or agents. The discount γ summarizes the construction horizon. Interface work includes choosing a statement at the right generality, selecting a memorable name, documenting it, controlling imports, and maintaining it across versions. These costs can dominate when the only scored outcome is today’s compilation. A short-horizon system may therefore be globally wasteful while locally optimal.

The principle yields discriminating predictions. Holding the theorem fixed, a shorter horizon should produce more local claims that are unused or single-use, less semantic investment in names, and shorter visible spans between definition and last reference. It does *not* predict that every machine claim is unreused, nor that every human claim is explanatory. Most importantly, it predicts an intervention: an AI given downstream library and maintenance objectives should cross the observed divide. Recent systems that explicitly extract reusable theorems or refactor generated proofs provide early evidence for exactly this mobility (Zhou et al., 2024; Fu et al., 2026; Lu et al., 2026; Lyu et al., 2026).

3 Paired observation at three resolutions

3.1 Corpus and estimand

We use the local NuminaMath-LEAN proof-artifact shards (AI-MO, 2025). A record is retained when human and prover proofs are both available and marked valid, the human ground-truth annotation is complete, and the UUID is unique. This gives 3,635 pairs from 12 coarse source groups. Each pair has the same formal statement; the proof bodies differ. We analyze only the final target declaration because 158 human artifacts contain helper declarations before it. Pooling those helpers would give one side an accidental miniature library.

“Human” denotes the dataset’s ground-truth proof provenance and “AI” its prover proof. It does not mean unaided human cognition versus an isolated model. Selection, shared libraries, model training data, and annotation workflows remain possible causes. Pairing removes theorem identity, not those process differences. Our estimand is consequently descriptive: how these two artifact tracks differ under the workflows that produced them.

At source level, a Unicode-aware parser removes comments and strings, locates named **have** claims in the final proof body, and counts later exact-name references. The counting window stops at same-name redeclaration, a conservative guard against shadowing. A claim has *long visible reach* when its last reference occurs after at least one intervening claim boundary. Names matching generic forms such as **h**, **h7**, **this**, **step2**, and **aux** are classified as placeholders. Results are pooled at claim level, with 95% intervals from resampling the 12 source groups. Pairwise proof metrics use Wilcoxon signed-rank tests; the effect sizes, source consistency, and sensitivity checks matter more than the very small p -values.

Lexical reference is not semantic use. To validate it we elaborate a preselected 312-pair sample against current Lean/Mathlib and inspect each **letE** in the final theorem’s proof term. A de-Brujin-aware traversal counts occurrences of a binder while increasing depth under nested binders; source claims are aligned to term binders by exact-name longest common subsequence in preorder. Of 624 tasks, 604 compile; 298 pairs have both sides. The alignment recovers 3,215 of 3,227 source claims in successful tasks (99.6%). A unique-name sensitivity gives the same direction. The 20 failures reflect toolchain drift and are reported, not silently replaced.

Finally, we score the 624 source documents in the preselected sample with a local Q4_K_M quantization of Goedel-Prover-V2-8B (Lin et al., 2025). Each token receives negative log probability given up to 8,192 preceding tokens. At each **have** boundary we compare the next 4, 8, 16, or 32




Figure 1: Three consumers of local claims. Source results use all 3,635 theorem pairs; term results use the 298 pairs that compiled on both sides. Error bars are 95% source-cluster intervals. “Descriptive” means a name outside the conservative placeholder list; “long reach” means the last explicit reference crosses at least one later claim boundary.

tokens with the preceding window and with eligible non-boundary locations in the same document. This is model-relative surprise under a prover trained on Lean-like text, not a measurement of human cognitive surprisal. Training overlap and affinity for the prover track are unexcluded. We therefore treat it as an exploratory timing assay rather than evidence of understanding.

3.2 Local claims have different fates

The paired proofs are not radically different in raw size: median final-body length is 126 whitespace tokens for the human track and 145 for AI, with a median within-pair increase of five tokens (source-cluster interval 2–15). They decompose that work differently. The median number of named claims is one versus four; pooled density is 1.97 human and 3.17 AI claims per 100 source tokens.

The important contrast is what happens after introduction (figure 1). Human claims receive 1.42 explicit references on average (source-cluster interval 1.32–1.56), compared with 0.86 for AI (0.83–0.88). Of human claims, 22.8% have zero visible uptake and 25.8% have multiple references. For AI the corresponding values are 36.4% and 12.0%. All 12 source groups show lower references per AI claim and higher zero-uptake share. Human claims also more often remain live across a later claim boundary: 47.6% versus 41.7%, in the human direction for 11 of 12 groups. The result is not simply that AI proofs are longer; it concerns the conditional fate of an introduced abstraction.

Names show an even larger difference. Placeholder forms account for 58.7% of human claims and 88.9% of AI claims. The pooled human name distribution has 3,598 types, Shannon entropy 8.52 bits, and effective vocabulary 367; the AI distribution has 1,686 types, 6.27 bits, and effective vocabulary 77 despite containing 75% more name tokens. Length is a crude but convergent check: 23.4% of human names and 9.4% of AI names have at least four characters. This is consistent with interface investment, but it is not by itself proof of semantic quality. A short conventional name may be excellent, and a long name may be decorative.

The elaborated term analysis both confirms and corrects the source story. Among aligned binders, 14.0% of human versus 23.4% of AI claims have zero occurrence in the final theorem term. Thus the AI track more often externalizes transient state that the certificate does not retain. But the rate of multiple term occurrences is essentially equal: 35.0% human and 35.9% AI. This is a substantive falsification of the simplest “AI does not reuse” narrative. Contextual tactics can duplicate an available fact many times during elaboration; raw mean use counts have enormous outliers and are not cognitively interpretable. What survives is more precise: AI introduces more




Figure 2: Source-group consistency. Each point is one of 12 source groups and shows AI minus human. References per claim are lower for AI in every group; zero-uptake is higher in every group; long visible reach is lower in 11 of 12. Point sizes encode paired sample count.

disposable source state, while kernel sharing among retained binders is not uniquely human.

3.3 One theorem, two temporal organizations

An illustrative pair makes the aggregate distinction legible. The theorem asks Lean to infer $x = 1$ from

$$2^{x^2-2x} + 3^{x^2-2x} + x^2 - 2x + \frac{1}{6} = 0.$$

The human-provenance script names the governing object $f(t) = 2^t + 3^t + t$, proves that f is strictly increasing, packages that fact as injectivity, and uses it to obtain $x^2 - 2x = -1$. Its three named claims are `e1`, `f_incr`, and `f_inj`; each is referenced once, and two remain live across a later boundary.

The AI script instead proves the key equality through two cases, below and above -1 , separately rederiving bounds for 2^t and 3^t . It introduces 24 claims named from `h1` through `h10`, with names reused in nested scopes. Eight aligned binders have zero final-term use. Both proofs check, and both recognize monotonic behavior. The difference is where that recognition is stored. The human version turns it into an addressable invariant; the AI version distributes it across case-local inequalities. One creates a small interface inside the proof, the other a trace of successful search.

This pair was selected because it is clear, not because it is representative by itself. Reverse examples exist. Its evidential role is to interpret the aggregate contrasts, whose paired and source-stratified directions were fixed independently.

3.4 Claim boundaries carry a short information pulse

Human conversational production regulates information in time: pauses and backchannels appear near changes in information rate (Bergey and DeDeo, 2024). If a named claim is an explanatory boundary rather than only syntax, it might likewise concentrate material that a predictive model finds locally unexpected.




Figure 3: Exploratory model-relative information around source `have` boundaries. Curves show negative log probability relative to the preceding 16 tokens; bands are proof-weighted 95% standard-error intervals. The model is Goedel-Prover-V2-8B in raw-code mode. The result may reflect training overlap, naming conventions, or prover-style affinity and is not a human surprisal measurement.

The exploratory result is compatible with that prediction (figure 3). Across the 174 pairs in which both proofs contain eligible boundaries, the median eight-token boundary excess— post-boundary minus pre-boundary NLL, corrected by same-document control positions—is 0.345 nats per token for human scripts and 0.024 for AI. The paired difference is -0.235 ($p = 3.9 \times 10^{-5}$). Removing the `have` keyword and claim name leaves an excess of 0.200 versus 0.042 ($p = .021$). The effect is present in 4- and 8-token windows and disappears by 16 and 32 tokens. Thus it is a local pulse, not sustained global difficulty.

The model assigns the AI documents much lower NLL overall (median 0.73 versus 1.15 nats/token), a warning that it is closer to AI proof style or data. We therefore do not infer that humans are “more creative.” The narrower observation is temporal: relative to controls within the same document, the human boundary marks a brief predictive discontinuity. Together with descriptive names and longer visible span, this suggests that some human claims package a change in explanatory state. A direct reader study and models screened for contamination are needed before calling the pulse cognitive.

4 From provenance to objective

The horizon principle would be weak if it merely renamed the observed label. Its main content is that the difference should move when the objective moves. Four recent lines of work support this prediction. REFACTOR learns to extract theorems from completed Metamath proofs; its 16 new lemmas are used hundreds of times on average after refactoring and improve downstream proving (Zhou et al., 2024). Proof-Refactor begins from the diagnosis that a compile-first objective produces proofs deficient in readability, modularity, maintenance, and reuse, then explicitly separates fragment

extraction, helper design, proof, and repair (Fu et al., 2026). Lean Refactor treats length, checking cost, and version robustness as distinct objectives rather than equating quality with compression (Lu et al., 2026). Rtl2lean asks a generator to build a four-layer theorem library and reports that 287 of 358 foundational lemmas are available for automatic reuse (Lyu et al., 2026).

These systems do not prove that their artifacts match mature human libraries, and their metrics are not directly comparable with ours. They establish the more fundamental point: machines can be made to externalize persistent mathematical memory. The observed human–AI difference is therefore best read as a difference between present production regimes. The causal experiment now suggested is straightforward: assign the same base model the same theorem stream under (i) per-theorem compile reward and (ii) downstream cost, maintenance, and human-retrieval reward. Measure dead binders, claim reach, API survival, later citation, repair cost, and reader comprehension. If the horizons do not separate, our account fails.

The distinction also clarifies DeDeo’s discussion of “alien lemmas.” Complexity theory implies that many truths have no short proof in a fixed system, while mathematical practice seems to select truths with good reasons (DeDeo, 2024b). An AI may discover a lemma that dramatically compresses later certificates yet remains unrecognizable to human mathematical ontology. Such a lemma has high machine amortized value and low human interface value. It is neither simply good nor bad. It exposes two horizons with different consumers. The correspondence problem then becomes an engineering and epistemic demand: construct a bridge that preserves the machine compression while letting human readers track why the abstraction matters (DeDeo, 2024a; DeDeo and Duede, 2026).

5 Cognitive implications: abstraction as memory allocation

The source-level difference suggests a cognitive interpretation more specific than “humans are abstract” and more defensible than reading mental states directly from code. Constructing a proof requires allocating limited representational capacity. A solver can preserve the past by keeping facts in the local context, by naming a relation that summarizes many facts, or by retrieving a library theorem that packages an older derivation. These are memory operations at different scales. The horizon determines which compression is worth performing.

At the scale of one search episode, an intermediate claim can function like a checkpoint. It freezes a successful subcomputation, constrains the remaining goal, and provides the next tactic with an enlarged context. Generic labels are adequate because identity can be recovered from position. Dead elaborated binders are not necessarily mistakes on this account. They are traces of a search process whose utility was exhausted before the final certificate. Removing them may improve the artifact while making the original generation harder to understand, just as deleting a scratch calculation produces a cleaner solution but erases evidence about discovery.

At a longer scale, an abstraction is a lossy compression chosen for transfer. The statement forgets derivational detail while preserving a relation expected to matter again. A descriptive name supplies a retrieval cue; a proof boundary asserts that the interior can be temporarily ignored; a general signature enlarges the class of future contexts in which the package applies. This compression is not free. Choosing the wrong invariant can increase working-memory demand, and exporting too many lemmas makes the library harder to search. Long-horizon construction is therefore not maximization of lemma count. It is selection of a small basis whose elements remain useful after the local context has vanished.

The equality of multi-use term binders is theoretically important here. Kernel sharing and communicative abstraction are different achievements. Automation can expand a generic hypothesis

into many term occurrences without ever identifying a concept for a reader. Conversely, a well-chosen explanatory lemma can deserve a name even when invoked once in the current theorem, because its intended reuse lies in another proof or another mind. Any cognitive account based only on the certificate DAG will confuse these cases. The units of human mathematical thought are not guaranteed to coincide with the compiler’s units of common-subexpression elimination.

The local surprisal pulse offers a tentative temporal signature of this selection. A useful boundary often arrives when prior manipulations have made a new description possible: monotonicity, invariance, symmetry, or a change of representation. The words immediately after the boundary can therefore be hard to predict from nearby syntax even though the resulting claim makes the rest of the argument easier. In conversational terms, the constructor briefly spends information to renegotiate common ground. The pulse’s disappearance after 16 tokens is consistent with a boundary event rather than generally ornate prose. But this interpretation must remain provisional until human reading time, recall, and transfer are measured directly.

This perspective also distinguishes discovery from presentation without separating them absolutely. A source proof can be a search trace, a pedagogical artifact, or both. Humans often refactor scratch work into a proof idea; present generators are commonly evaluated before that second pass. The central divergence may thus be procedural: one workflow includes consolidation, naming, and anticipation of reuse, while another terminates at verification. A cognitively serious comparison should align stages—first attempts with first attempts, refactored artifacts with refactored artifacts—and observe how representations change between them. The horizon principle predicts that the largest human–AI gap will occur not in local inference, where both can exploit formal context, but in deciding what part of a successful episode should become shared memory.

6 Amortized proof complexity

Ordinary proof complexity asks for the size of a certificate for one formula in a fixed proof system. This leaves library construction outside the objective. Let $\Phi = (\varphi_1, \dots, \varphi_T)$ be a stream of targets. For a checked auxiliary library L and proofs P_t that may use it, define

$$\text{APC}_\lambda(\Phi) = \min_{L, P_1, \dots, P_T} \left[\lambda |L| + \sum_{t=1}^T |P_t \mid L| \right], \quad (2)$$

subject to every declaration in L and every P_t checking in the base environment. The factor λ prices the design and maintenance of shared structure. A richer version can price names, dependency fragility, retrieval, and human validation separately. An online version charges when each declaration is introduced and compares cumulative cost with the best hindsight library.

Equation (2) changes what counts as an optimal proof. A declaration that lengthens the first episode may lower total cost if it compresses later proofs. A generic helper that is short but impossible to retrieve may fail to pay back its interface cost. A benchmark of independent theorems with a reset context implicitly sets future reuse to zero and selects a compile-first policy. Even pass@ k with enormous search budgets remains a measure of finding current certificates, not of learning a mathematical language in which later certificates become cheap.

The connection to circuit sharing and extension variables is instructive but limited. Naming a repeated subexpression turns tree accounting into DAG accounting; introducing a reusable theorem resembles adding a subroutine or grammar rule. Yet Lean claims are dependent types, tactics alter elaboration, definitions compute, and public theorem statements carry semantic and maintenance costs. We do not claim an equivalence with Extended Frege. The robust lesson is accounting level: single-certificate complexity cannot see investments whose return lies across certificates.

This produces new empirical questions at the boundary of complexity and cognition. How large can online library regret be for a bounded-context agent? Which theorem orderings permit discovery of the hindsight-optimal abstractions? Does a human-readable interface impose only polynomial overhead on a machine-efficient lemma basis, or can correspondence itself be exponentially costly? DeDeo’s “hard proofs” dilemma suggests that mathematical culture may be understood as a selection process over theorem streams and representations, concentrating attention where useful proofs and reasons coincide (DeDeo, 2024b). AI changes both the search distribution and the consumers for whom that coincidence must hold.

7 What the result does and does not establish

The evidence supports a narrow, structural claim. In this paired corpus, the AI track uses more source-level local state; more of it is visibly unreferenced, more of it vanishes from the final term, its names are more generic, and its visible temporal reach is shorter. It does not support a claim that machine proof terms lack sharing: retained binders are multiply used at the same rate. Nor does it show that human artifacts are globally reusable, since our primary analysis ends at the target declaration. The distinction between episode and persistent interface remains to be tested on theorem sequences.

Several limitations are consequential. First, source provenance is observational. Second, the lexical parser recognizes named `have`, not every `let`, suffices, anonymous claim, or tactic-generated subgoal. Third, compilation leaves 298 complete pairs in the term sample; although their source patterns resemble the full corpus, toolchain attrition can select on style. Fourth, binder occurrence is syntactic and can be amplified massively by automation. Fifth, the surprisal model is a quantized theorem prover run without a chat template and may have seen related text. Its boundary effect is exploratory. Sixth, meaningful names and long spans are only proxies for explanatory value. Interactive links from written proof steps to Lean state have improved comprehension in a small reader study, but our boundary measures still require a direct experiment (Kambhamettu et al., 2026).

The representation audit imposes a further correction. The earlier extractor interned raw `Expr.bvar` 0 nodes globally, merging variables belonging to unrelated binders. This can create false reuse paths in an unexpanded term DAG. The extractor now instantiates loose indices with scope-specific free-variable identities, and a regression test constructs two lambdas and requires two distinct variable nodes. Declaration graphs are unaffected, but earlier semantic interpretations of root-term topology should be withdrawn pending scope-correct reanalysis. The binder-use results reported here use a separate depth-aware traversal and do not inherit that flaw.

Three observations would falsify or sharply narrow the horizon account. The claim fails if an objective-matched human–AI experiment retains the same difference; if downstream reward does not produce longer-lived machine interfaces; or if the proposed interface measures fail to predict future proof cost, maintenance, or reader comprehension. Conversely, a machine system that builds a durable, intelligible library is not an exception to be explained away. It is the decisive prediction.

7.1 Three decisive experiments

The first experiment is a horizon swap. Give humans and machines the same ordered Lean tasks and the same base library. In one condition, discard each proof after checking it. In the other, make every later task cheaper when an earlier abstraction is reusable and charge for broken or unretrievable APIs. The critical statistic is the interaction between constructor and condition. A

stable provenance effect would contradict the objective-level account; convergence within condition would support it.

The second is a consumer swap. Hold certificates fixed and ask three consumers to operate on alternative refactorings: the kernel, a new prover with bounded context, and a human reader. Kernel cost measures certificate function; downstream search measures machine interface value; delayed recall, error localization, and transfer measure human interface value. This design can reveal alien lemmas rather than assuming them: an abstraction has high value for one consumer and low value for another. A correspondence layer succeeds when it raises human value without destroying machine savings.

The third is a longitudinal library experiment. Allow an agent to choose which intermediate facts to export before seeing future targets, then measure cumulative cost and survival over theorem orderings and Lean versions. Compare its online choices with a hindsight refactoring of the whole stream. This directly estimates library regret and separates genuine anticipation from post hoc compression. It also makes failure scientifically useful. A system that proves every target but accumulates high regret has mastered deduction without mastering the allocation of mathematical memory.

Together these experiments move the field from artifact classification to causal cognitive science. They ask not whether one population’s code looks more elegant, but which incentives produce representations that remain valuable when context, consumer, or time changes. The unit of analysis is no longer a proof in isolation. It is the transformation of successful reasoning into shared capacity.

8 Conclusion: a proof is a wager on its next use

A checked certificate has no intrinsic audience beyond the kernel. Proof construction becomes mathematical practice when it anticipates another consumer: a later step, a future theorem, a maintainer, a reader, or another agent. Human- and AI-provenance Lean proofs of the same statement currently diverge most clearly at this point. One track more often packages facts as named, longer-lived hand-offs; the other more often records disposable local state. After elaboration, both can exhibit extensive sharing. The difference is not deductive validity and need not be a fixed property of the constructor. It is the future for which the proof was optimized.

This view turns proof quality from a static aesthetic judgment into an amortization question. Who will use this boundary, over what distance, and what work will it save? It also identifies a blind spot in current evaluation: a sequence of isolated successes can grow indefinitely without becoming a library. The next generation of AI-for-mathematics benchmarks should therefore measure not only whether a system can prove today’s theorem, but whether today’s proof changes the cost and intelligibility of tomorrow’s.

References

- AI-MO. NuminaMath-LEAN: A large-scale dataset of competition problems formalized in Lean 4. Hugging Face dataset, 2025.
- Claire Augusta Bergey and Simon DeDeo. From “um” to “yeah”: Producing, predicting, and regulating information flow in human conversation. 2024. arXiv:2403.08890.
- Katherine M. Collins, Simon Frieder, Jonas Bayer, Jacob Loader, Jeck Lim, Peiyang Song, Fabian Zaiser, Lexin Zhou, Shanda Li, Sam Looi, Joshua B. Tenenbaum, Umang Bhatt, Adrian Weller, Jose Hernandez-Orallo, Cameron E. Freer, Valerie Chen, and Ilia Sucholutsky. Characterizing initial human–ai proof formalization workflows, 2026. arXiv:2606.04273.
- Simon DeDeo. Alephzero and mathematical experience. *Bulletin of the American Mathematical Society*, 61(3):375–386, 2024a. doi: 10.1090/bull/1824.
- Simon DeDeo. Hard proofs and good reasons, 2024b. arXiv:2410.18994.
- Simon DeDeo and Eamon Duede. A correspondence problem for mathematical proof, 2026. arXiv:2603.13680.
- Yiming Fu, Peixuan Liu, Zichen Wang, and Kun Yuan. Proof-refactor: Refactoring generated formal proofs into modular artifacts, 2026. arXiv:2606.03743.
- Hita Kambhamettu, Will Crichton, Sean Welleck, Harrison Goldstein, and Andrew Head. Making written theorems explorable by grounding them in formal representations, 2026. arXiv:2604.02598.
- Yong Lin, Shange Tang, Bohan Lyu, Ziran Yang, Jui-Hui Chung, Haoyu Zhao, Lai Jiang, Yihan Geng, Jiawei Ge, Jingruo Sun, Jiayun Wu, Jiri Gesi, Ximing Lu, David Acuna, Kaiyu Yang, Hongzhou Lin, Yejin Choi, Danqi Chen, Sanjeev Arora, and Chi Jin. Goedel-prover-v2: Scaling formal theorem proving with scaffolded data synthesis and self-correction, 2025. arXiv:2508.03613.
- Jialin Lu, Soonho Kong, Rodrigo Stehling, Kaiyu Yang, Zhangyang Wang, Weiran Sun, and Wuyang Chen. Lean refactor: Multi-objective controllable proof optimization via agentic strategy search, 2026. arXiv:2605.20244.
- Hongqin Lyu, Junxing Dong, Yonghao Wang, Zhiteng Chao, Tiancheng Wang, and Huawei Li. Rtl2lean: Automated RTL-to-Lean translation with hierarchical theorem generation and lemma reuse, 2026. arXiv:2607.16855.
- Yehuda Rav. Why do we prove theorems? *Philosophia Mathematica*, 7(1):5–41, 1999.
- Scott Viteri and Simon DeDeo. Epistemic phase transitions in mathematical proofs. *Cognition*, 225: 105120, 2022. doi: 10.1016/j.cognition.2022.105120.
- Jin Peng Zhou, Yuhuai Wu, Qiyang Li, and Roger Grosse. REFACTOR: Learning to extract theorems from proofs, 2024. arXiv:2402.17032.

A Reproducibility and robustness

All reported paired source results are produced by `code/paired_horizon.py`; scope-aware term use by `code/extract_binder_use.py` and `code/analyze_binder_use.py`; and surprisal by `code/prepare_surprisal.py`, `code/token_surprisal.cpp`, and `code/analyze_surprisal.py`. Machine-readable summaries live in `results/horizon/`. Seeds and inclusion rules are serialized with the estimates.

Source intervals resample 12 source groups with replacement and recompute pooled claim ratios. Every source group has lower explicit uses per AI claim and higher AI zero-uptake; 11 of 12 have lower AI long-reach share. The term alignment’s unique-name subset contains 753 human and 1,258 AI claims; zero-use shares are 13.8% and 23.8%, respectively. The surprisal boundary contrast is present for four- and eight-token windows, marginal at 16, and absent at 32. Excluding the keyword and claim name preserves the four- and eight-token effects.

The source direction also survives a coarse automation sensitivity. For each of `linarith`, `nlinarith`, `norm_num`, `omega`, `ring`, and `simp`, we separately retain pairs in which both sides invoke the tactic and pairs in which neither does. Across all 12 strata, AI claims have fewer explicit references per claim and a higher zero-uptake share. This does not equate full tactic workflows, but rules out any one of these common families as a sufficient explanation.

B Why the old network result is not the new result

The prior project replicated Viteri and DeDeo’s fixed-tail power-law results and modeled epistemic phase transitions across Coq and Lean. Those are properties of declared network constructions. They do not supply a representation-invariant authorship signature. Tail exponents move when the minimum degree or expansion boundary moves; human–AI directions can reverse under proof-value normalization; and modeled theorem belief is close to saturation in both tracks.

The scope audit strengthens the reason for restraint. A de Bruijn index identifies a binder by distance in its local expression, not by a global name. Interning equal raw indices across scopes turns distinct variables into one graph node. The corrected traversal carries a scope environment and replaces each entered binder with a fresh free-variable identity before hashing expressions. The regression proof has two independent identity lambdas; the old representation produces one shared variable leaf, while the corrected representation produces two. None of the main analyses in this paper uses a raw-term reuse tail.

C A proposed causal benchmark

A strong test would sample theorem sequences, not independent goals. Randomize capable agents to a compile-first condition or a library condition. Both receive the same base library, target order, Lean version, and inference budget. The library condition is charged for current proof cost but rewarded for reductions in later search, version repair, and blinded reader retrieval. Public declarations must be generated before future targets are revealed to prevent hindsight leakage.

Primary outcomes should include cumulative checked-token or search cost, online regret against a hindsight refactoring, survival of generated APIs under version changes, cross-target citations, dead local binders, and delayed human comprehension. A swap condition—humans under disposable single-goal incentives and AI under persistent-library incentives—would separate constructor from horizon. The strongest result for the present theory would be a crossover interaction rather than a persistent main effect of provenance.