

Proofs for Now and Proofs for Later

The Amortization Horizon of Human and Artificial Reasoning in Lean

Working paper

2 August 2026

Abstract

A checked proof establishes that a proposition follows, but not what future it was built to serve. We propose the *amortization horizon*: the range over which a constructor expects an abstraction to repay its cost. A proof produces a certificate, working memory, and an interface for users. In 3,630 pairs of Lean proofs of identical statements, AI scripts create more claims, but their absolute supply advantage reverses between one and two later references; they are also about one-third as likely to state a callable family and more generically named. Elaboration then inverts the reuse story: AI boundaries more often vanish, but surviving AI binders more often fan out to multiple term occurrences. This consolidation–composition split distinguishes deliberate source reference from compiler fan-out. Certificate topology also fails consumer-invariance: Lean versions encode the same source claim differently, while extensionally equivalent representatives can differ exponentially in tree size. An exploratory model finds an information pulse at human claim boundaries, attenuated at AI boundaries. The divergence lies not in deductive power or biological kind, but in which future consumer construction serves. We formalize the principle as library-amortized proof complexity over theorem streams, a quantity single-theorem benchmarks cannot observe.

1 The missing future of a checked proof

A theorem prover settles a sharply posed question: did this term check at this type? That question is essential, but it is not the whole function of proof. A mathematician also uses a proof to find the invariant that makes a result intelligible, to remember a route through an argument, to expose a lemma that will solve the next theorem, and to leave an object another person can repair (Rav, 1999). These uses extend beyond the instant of verification. Two certificates of the same statement may thus agree perfectly as kernel inputs while serving different futures.

This paper argues that this future is the most useful level at which to compare present human- and AI-provenance Lean proofs. The common contrast—human proof versus machine proof—encourages an essentialist question: what signature does the identity of the constructor leave in a proof graph? That question is poorly controlled. Theorem, library, tactic, model, and representation all change graph statistics. More importantly, it mistakes an objective for an organism. Most generated proofs are rewarded when the current declaration compiles. Much human formalization occurs inside a longer activity of exposition, maintenance, and library construction; early workflow evidence also finds that users want AI assistance while preserving high-level control (Collins et al., 2026). In resource-rational terms, different artifacts may be optimal under different costs even before any difference in cognitive machinery is posited (Lieder and Griffiths, 2020).

We call the relevant variable the *amortization horizon*: how far into future reasoning the constructor expects a present abstraction to pay. At a short horizon, a local claim can be useful merely because it records a search result for the next tactic. Its name need only be distinct; its statement need not be reusable; it may disappear under elaboration. At a long horizon, names, boundaries, generality, and stable dependencies become investments. Their up-front cost can be recovered across later steps, theorems, versions, and readers.

This account makes four contributions. First, it separates three products of formal construction: certificate, episode structure, and interface. Second, it tests their separation in 3,630 exact- statement human–AI pairs, with a scope-aware decoding of the dataset’s production proof terms and a two-version representation audit. Third, it offers an exploratory information-theoretic test inspired by work on how speakers regulate surprise in conversation (Bergey and DeDeo, 2024). Finally, it proposes a sequence-level object, amortized proof complexity, that charges for a library and the stream of proofs it enables. This reframes a central evaluation problem: success at one theorem measures certificate production, not mathematical accumulation.

The argument deliberately narrows several earlier claims in this repository. Viteri and DeDeo’s proof-network work establishes that proof organization can support robust belief and reveals heavy-tailed reuse (Viteri and DeDeo, 2022). Our replication found those results informative but highly sensitive to the graph boundary. Near-saturated belief did not distinguish paired certificates, and a newly discovered de Bruijn-scope conflation makes the inherited unexpanded-term topology unsuitable as a semantic reuse measure. Rather than force authorship into those statistics, we use the failure to identify the missing object: who or what must use the proof next?

2 A proof has three consumers

The word “proof” compresses at least three artifacts.

Certificate. The kernel consumes an elaborated term and checks that it inhabits the target type. The certificate may contain local `let` binders and shared subexpressions, but the kernel does not require human-readable names, explanatory order, or a useful public API.

Episode structure. The constructor and an immediate reader consume the source script during the current attempt. A named `have` makes a fact addressable, reduces the active goal, and records search state. It can perform this role even if no later source token names it: a contextual tactic such as `linarith` may consume it implicitly, or elaboration may erase it. In cognitive terms this is a distributed representation: internal and external state jointly carry the task, and writing a claim can be an epistemic action that lowers later computation (Zhang and Norman, 1994; Kirsh and Maglio, 1994).

Interface. Future proofs and maintainers consume exported declarations, stable names, general statements, and explanatory boundaries. An interface is not just a set of dependencies. It is a retrieval scheme: it partitions a derivation into units that can be recognized and invoked without reconstructing their interiors. This audience-relative view accords with the dialogical conception of proof as discourse directed toward explanatory persuasion (Dutilh Novaes, 2018). DeDeo and Duede call attention to the associated correspondence problem: a formal derivation can exist while failing to track the proof idea that is supposed to justify its relevance (DeDeo and Duede, 2026). The interface can also certify something about its constructor: explanatory choices provide evidence of understanding that bare correctness does not. Cheap generation can sever this costly-signal function even while improving verification (Wojtowicz and DeDeo, 2025). For this meta-consumer, construction cost is not mere overhead but part of the evidence; automation may preserve the visible artifact while changing what its production warrants.

These consumers define different notions of reuse. A proof-term binder used twice is a sharing node in a directed acyclic graph. A source claim referenced after several later claims is working memory with temporal reach. An exported theorem cited by later files is persistent memory. None is a universal measure of modularity. In particular, raw occurrence counts can mistake tactic expansion for conceptual importance, while a theorem that is used only once may still identify the right explanatory cut. This instantiates Larkin and Simon’s distinction between informational and computational equivalence: the same content can impose different search and recognition costs on different operators (Larkin and Simon, 1987).

This separation imposes a methodological constraint. Let $R_c(e)$ be consumer c ’s response to artifact e , and write $e \sim_c e'$ when $R_c(e) = R_c(e')$. The *consumer-invariance principle* says that a statistic offered as a functional property for c must be constant on each $\sim_c$ class. For correctness, Lean’s kernel accepts any term of the target type; proof irrelevance makes proofs of one proposition definitionally interchangeable, and zeta reduction replaces a `let`-bound variable by its value (Lean Developers, 2026). Consequently, binder count, outdegree, and even proof-tree size are not intrinsic properties of the established judgment. They can still diagnose checking cost, an elaborator snapshot, or a consumer that reads that representation, but the observable and representative must be named. There is a simple impossibility corollary. Fix a proposition and let the consumer’s only response be kernel acceptance. All accepted proof terms then occupy one response class, so every consumer-invariant functional statistic is constant across them. Any nonconstant human–AI certificate statistic necessarily describes a richer response—time, memory, elaboration, inspection, retrieval, or provenance—rather than correctness itself. This is the proof-theoretic identity problem in empirical form: normalization and generality induce different

quotients over derivations (Došen, 2004), and an authorship statistic must declare which quotient, if any, it survives.

The failure can be exponential, not cosmetic. Let e_0 be the natural-number literal zero and $e_{n+1} = \text{let } x := e_n; x + x$. Retaining sharing gives $1 + 7n$ syntax-tree nodes in our Lean encoding; full zeta reduction gives $4 \cdot 2^n - 3$. At $n = 14$, Lean’s own reducer changes 99 nodes and 14 binders into 65,533 nodes and none, without changing what conversion can establish. Thus a proof graph is a representative of checked evidence, not a representation-free map of reasoning. Elaboration is better treated as a two-stage channel: first ask whether a visible source boundary is retained at all, then, conditional on retention, whether it is used zero, once, or many times. The empirical toolchain comparison below shows why those stages cannot be collapsed.

For an abstraction z and consumer c , a minimal decision model is

$$V_{\gamma,c}(z) = \sum_{u \in U_c(z)} \gamma^{d(z,u)} s_c(u) - c_{\text{make}}(z) - c_{\text{interface},c}(z), \quad (1)$$

where $U_c(z)$ is the set of anticipated uses, $s_c(u)$ is work saved at use u , and d is distance in steps, theorems, time, or agents. The discount γ summarizes the construction horizon. Interface work includes choosing a statement at the right generality, selecting a memorable name, documenting it, controlling imports, and maintaining it across versions. These costs can dominate when the only scored outcome is today’s compilation. A short-horizon system may therefore be globally wasteful while locally optimal.

The consumer index matters as much as the discount. The same lemma can have positive value for a retrieval-equipped prover and negative value for a reader who cannot recognize or trust it. There is consequently no consumer-free scalar called “lemma quality.” Horizon specifies when value must be recovered; audience specifies whose future work counts.

This is also a rate–distortion problem. Let $E \sim p(e)$ denote completed proof episodes and Z the interfaces retained from them. For consumer c , define $D_c(E, Z)$ as the expected extra work, error, or loss of correspondence when future tasks receive Z rather than the full episode. An interface policy trades representational rate against discounted distortion, for example by minimizing $I(E; Z) + \beta \mathbb{E}_\gamma[D_c(E, Z)]$ (Shannon, 1959). Consumer-invariance identifies the zero-distortion quotient; horizon and counterfactual width determine the future-task distribution. The same derivation can therefore have different optimal compressions for a kernel, a retrieval-equipped prover, and a person tracking the proof idea.

Distance is only one dimension of the wager. A claim can also range over counterfactual cases that the current derivation never visits. A parameterized lemma increases this *counterfactual width*: it packages a family of related facts rather than only the instance presently needed. Horizon measures how far an abstraction travels; width measures how many varying contexts it can enter. This echoes the distinction between instance-based and generic explanation (Linnebo, 2022). Both enlarge an abstraction’s amortization domain, and both cost something before they save work.

The principle yields discriminating predictions. Holding the theorem fixed, a shorter horizon should produce more local claims that are unused or single-use, less semantic investment in names, and shorter visible spans between definition and last reference. It does *not* predict that every machine claim is unreused, nor that every human claim is explanatory. Most importantly, it predicts an intervention: an AI given downstream library and maintenance objectives should cross the observed divide. Recent systems that explicitly extract reusable theorems or refactor generated proofs provide early evidence for exactly this mobility (Zhou et al., 2024; Fu et al., 2026; Lu et al., 2026; Lyu et al., 2026).

3 Paired observation at three resolutions

3.1 Corpus and estimand

We use the local NuminaMath-LEAN proof-artifact shards (AI-MO, 2025). A record is retained when human and prover proofs are both available and marked valid, the human ground-truth annotation is complete, and the UUID is unique. We then require both serialized artifacts to contain the same normalized final declaration signature. This audit removes one nominally valid prover artifact with no target and four mismatches: one wrong theorem join, one helper-only output, and two altered propositions. The result is 3,630 exact-statement pairs from 12 coarse source groups. We analyze only the final target declaration. In 158 pairs, both tracks contain identical pre-target declarations after whitespace and comment normalization; 11 additional AI artifacts

introduce helpers named only `lemma_1` or `lemma_2`, seven of which the target cites. Pooling these declarations would therefore mix shared problem scaffolding and side-specific output with the target-proof comparison, not reveal a human-only miniature library.

The dataset card defines its formal ground truth as a proof written by a human annotator and its prover field as a model proof produced during reinforcement-learning training (AI-MO, 2025). “Human” and “AI” below denote those artifact fields. This does not mean unaided human cognition versus an isolated model. Selection, shared libraries, model training data, and annotation workflows remain possible causes. Pairing removes theorem identity, not those process differences. Our estimand is consequently descriptive: how these two artifact tracks differ under the workflows that produced them.

At source level, a Unicode-aware scanner removes comments and strings and locates named **have** headers in the final proof body. Lean’s own parser then supplies the end of the complete construction tactic, so occurrences in a claim’s type or proof term cannot be mistaken for downstream use. It exactly locates 39,646 of 40,626 scanner candidates (97.6%); all primary source estimates use those claims. We count later exact-name references. The window stops at the tail of the nearest enclosing tactic sequence, or earlier at same-name redeclaration. A claim has *long visible reach* when its last reference occurs after at least one intervening claim boundary. Names matching generic forms such as `h`, `h7`, `this`, `step2`, and `aux` are classified as placeholders. A family claim has either binders before its type or a type beginning with `forall`/ $\forall$ (up to parentheses), so equivalent Lean spellings enter the same class. Because boundary count depends on how densely a script is decomposed, we separate adoption from duration and also report token distance and the fraction of available later boundaries crossed. Results are pooled at claim level, with 95% intervals from resampling the 12 source groups. Pairwise proof metrics use Wilcoxon signed-rank tests; the effect sizes, source consistency, and sensitivity checks matter more than the very small p -values.

Lexical reference is not semantic use. We therefore elaborate all 7,260 target artifacts against the dataset’s production snapshot, Lean/Mathlib 4.15.0. That version represents a source **have** as an application of `letFun` to a named lambda; Lean 4.33 instead uses a nondependent `letE`. A node-kind-only extractor would consequently report near-total erasure under one version. Our stack-safe decoder recognizes both semantic encodings. A de-Bruijn-aware traversal counts binder occurrences while increasing depth under nested binders, and exact-name longest common subsequence in preorder aligns source claims to decoded binders. A stack-safe dynamic program preserves full tree-occurrence counts while memoizing shared subexpressions; every status and code/toolchain hash is serialized. All 7,260 tasks complete.

As a representation sensitivity, we independently retain the preselected 312-pair scope-audit sample under Lean/Mathlib 4.33.0-rc1. Of 624 tasks, 604 compile and 298 pairs have both sides. Its 20 failures measure version attrition rather than production-time invalidity. Cross-version claims are compared only where both tasks compile and the source binder aligns under both decoders.

Finally, we score the 624 source documents in the preselected sample with a local Q4_K_M quantization of Goedel-Prover-V2-8B (Lin et al., 2025). Each token receives negative log probability given up to 8,192 preceding tokens. At each **have** boundary we compare the next 4, 8, 16, or 32 tokens with the preceding window and with eligible non-boundary locations in the same document. This is model-relative surprise under a prover trained on Lean-like text, not a measurement of human cognitive surprisal. Training overlap and affinity for the prover track are unexcluded. We therefore treat it as an exploratory timing assay rather than evidence of understanding.

3.2 Local claims have different fates

The paired proofs are not radically different in raw size: median final-body length is 125 whitespace tokens for the human track and 145 for AI, with a median within-pair increase of five tokens (source-cluster interval 2–15). They decompose that work differently. The median number of named claims is one versus four; pooled density is 1.96 human and 3.10 AI claims per 100 source tokens.

The Kimina-Prover report treats abundant **have** statements as evidence of “human-like” clarity (Numina and Kimi Teams, 2025). Yet count alone is non-identifying: any proof can be prefixed by arbitrarily many unused claims of `True` without changing its target. The empirical denominator matters too. Per 100 tokens, AI emits more claims adopted at least once (1.96 versus 1.52), but humans emit more claims that are multiply referenced (.470 versus .318), descriptively named (.805 versus .348), or generalized (.103 versus .052); every difference has a source-cluster interval excluding zero. The absolute advantage reverses between the first and second exact-name reference, and every source group favors humans at two or more. This *lexical consolidation*

crossover shows that decomposition—making a boundary—is not consolidation—making it addressable. A conservative site/DAG audit makes absolute multi-site supply denominator-dependent, but at equal node count human graphs are denser and more branching; maximum chain length is unresolved (appendix). Thus the crossover identifies repeated textual addressability, not distinct acts of use by a reader.

The important contrast is what happens after introduction (figure 1). Human claims receive 1.35 explicit references on average (source-cluster interval 1.26–1.49), compared with 0.82 for AI (0.79–0.85). Of human claims, 22.2% have zero visible uptake and 24.0% have multiple references. For AI the corresponding values are 36.9% and 10.3%. All 12 source groups show lower references per AI claim and higher zero-uptake share. Human claims also more often remain live across a later claim boundary: 42.4% versus 35.7%, in the human direction for all 12 groups. The result is not simply that AI proofs are longer; it concerns the conditional fate of an introduced abstraction. In 834 pairs within 10% in body length, uses remain 1.25 versus 0.98, zero uptake 24.5% versus 34.3%, and long reach 42.5% versus 36.9%. Reach has an extensive and an intensive margin. Human claims are explicitly adopted 77.8% of the time versus 63.1% for AI; conditional on adoption, they receive 1.74 versus 1.30 explicit calls, and their last reference lies 35.0 versus 23.9 source tokens beyond the completed construction. Counts of intervening claims are endogenous to decomposition density: given adoption and at least one available later boundary, the probability of crossing one is indistinguishable (83.6% versus 85.0%). The fraction of available boundaries crossed is likewise unresolved (63.7% versus 62.0%). We therefore interpret the aggregate difference as lower AI adoption, not a representation-free lifetime constant. Indeed, within 834 length-matched pairs neither conditional-duration measure separates the tracks; in 206 pairs with the same positive claim count, normalized duration is nearly identical (64.3% versus 64.2%). The robust source divergence is explicit adoption, not how long an adopted claim lives. The result is not arithmetic from claim count: in 206 pairs with equal positive counts, uses remain 1.29 versus 1.04 (AI-minus-human interval -0.411 to -0.157) and zero uptake 21.9% versus 30.0% (0.053 to 0.109). Among names declared only once per proof, uses are 1.46 versus .88, zero uptake 19.7% versus 34.0%, and long reach 46.8% versus 40.1%.

The strongest direct signature of counterfactual width is grammatical. A local lemma can turn one fact into a callable family either with binders—`have sum_range (n : Nat) : ...`—or with an explicitly universal type, `have sum_range : forall n, ...`. Counting only the first spelling produces a dramatic but non-invariant fourteenfold contrast. Quotienting the two forms still leaves a large difference: 761 of 14,519 human claims (5.24%) and 422 of 25,127 AI claims (1.68%) state a family, with the human direction in all 12 source groups. Within identical-theorem pairs, a generalized claim occurs only on the human side in 310 cases, only on the AI side in 102, and on both in 145. This remains a syntactic classifier, not proof that a statement chose the best generality. But it records a real construction cost: making an intermediate result vary over an argument. Conditional on paying that cost, AI families are not shorter-lived by this pooled measure: 63.2% of the human and 67.5% of the AI generalized claims remain live across a later claim boundary. Here too the difference lies in how often the abstraction is selected, not an inability to use it once made. The classifier has the same functional correlate in both tracks. Minimum-cost matching to nearby instance claims, with a quarter-proof position caliper, retains 248 human and 199 AI proofs: family claims are respectively 15.0 and 20.6 points more likely to be adopted, and 26.8 and 37.0 points more likely to be multiply referenced; all four intervals exclude zero. Among 79 pairs eligible on both tracks, the multi-reference uplift is larger for AI (interval 3.5–23.2 points). The association is not explained by families occurring earlier. It remains descriptive, but the rarer AI families are especially active when selected. In the 206 pairs with the same positive claim count, family shares remain 4.36% versus 2.56%. The difference is not carried by closed-goal computation: after excluding 416 pairs in which either side invokes `native_decide`, family shares remain 5.11% versus 1.72%.

Names show an even larger difference. Placeholder forms account for 58.9% of human claims and 88.8% of AI claims. The pooled human name distribution has 3,578 types, Shannon entropy 8.53 bits, and effective vocabulary 369; the AI distribution has 1,662 types, 6.28 bits, and effective vocabulary 78 despite containing 72% more name tokens. Length is a crude but convergent check: 23.8% of human names and 9.4% of AI names have at least four characters. This is consistent with interface investment, but it is not by itself proof of semantic quality. A short conventional name may be excellent, and a long name may be decorative. A paired within-proof embedding assay finds no robust human advantage after chance and claim-count normalization (appendix), so we make no semantic naming claim. Nor do the proxies collapse into one style score. Across the 1,816 pairs with claims on both sides, Spearman correlations among proof-level human advantages in uptake, reach, family construction, and naming range only from .04 to .34. We therefore keep the coordinates separate.

The elaborated term analysis both confirms and corrects the source story. All 7,260 production- snapshot tasks compile; exact-name alignment retains 14,469 of 14,519 human and 25,090 of 25,127 AI claims. Among




Figure 1: Three consumers of local claims. Panel A gives pooled source-claim supply at each minimum later exact-name-reference count (log scale), crossing at two. Source panels use all 3,630 pairs; the term panel uses complete pairs. Error bars in B–D are 95% source-cluster intervals. “Non-generic” excludes placeholder forms; “local family” unions binder and universal spellings.

aligned binders, 9.8% human versus 21.8% AI have zero occurrence in the final term. But multiple occurrence is 32.2% versus 34.2% (AI-minus-human interval -0.04 to 5.08 points). Conditional on occurring at least once, however, multi-use is 35.7% versus 43.7%, higher for AI in 11 of 12 groups (difference interval 4.8 – 12.7 points). The sharper result is a polarized source-to-certificate channel: zero-or-multiple occurrence is 42.0% human and 56.0% AI (difference 14.0 points, interval 9.9 – 20.7), leaving one-to-one mapping at 58.0% versus 44.0%. Within both the source-single-use and source-multiple claims, AI claims are more likely to disappear and more likely to be duplicated; polarization is higher for AI in all 12 matched tactic strata. Contextual tactics can amplify a fact massively, so raw means are not cognitive reuse measures.

The inversion holds claim by claim. Among source-single-use claims, 20.4% human versus 31.1% AI become multi-occurrence binders; among source-multiple claims, 58.4% versus 67.0%. AI disappearance is also higher in both strata. Source reference and term fan-out are distinct channels.

Family syntax nevertheless predicts compiled function. Repeating the one-to-one position match on aligned binders (quarter-proof caliper; 245 human and 199 AI proofs), family-minus-instance disappearance is -2.3 points human (interval -5.3 to -1.0) and -8.4 AI (-10.5 to -5.6). The corresponding multi-use uplift is 17.7 points human (10.4 – 25.4), but -0.3 AI (-9.9 to 16.4). The direct AI-minus-human interaction among 79 pairs eligible on both tracks remains unresolved (interval -19.0 to 7.1). Thus families preferentially survive on both tracks; the different composition estimates are suggestive, not an established provenance difference. This validates the classifier functionally without making occurrence count a general cognitive-reuse measure.

The three resolutions reveal a *consolidation–composition split*. The human track more often turns intermediate state into a visibly adopted, repeatedly referenced, named, and general interface; the AI track more often leaves its bound fact absent from or amplified within the returned certificate. This is a surface–certificate inversion: explicit source reference is a construction decision, whereas term fan-out is partly an effect of tactics and elaboration. Aggregate term multi-use does not favor humans, and AI families are not detectably shorter-reaching in the pooled source comparison. The present divergence is therefore better localized to consolidation policy—where reusable compositional structure is made explicit—than to deductive capacity.

3.3 One theorem, two amortization horizons

An extreme pair makes the distinction legible. Define an integer sequence by $f(0) = 2001$, $f(1) = 2002$, $f(2) = 2003$, and $f(n+3) = f(n) + f(n+1) - f(n+2)$; the target asks only for $f(2003) = 0$. The human-provenance script proves by induction a simultaneous family

$$f(2k+2) = 2003 + 2k, \quad f(2k+1) = 2002 - 2k, \quad f(2k) = 2001 + 2k,$$

names it `f_formula`, and instantiates $k = 1001$. The AI script is essentially one command: `norm_num [f]`. One proof invests in a counterfactually wide description; the other asks the elaborator to settle the present instance.

Both check, but their representation costs reveal why certificate topology needs a named consumer. The human term has 278,395 nodes under full tree-occurrence accounting. Repeated expansion of the AI term




Figure 2: Source-group consistency. Each point is one of 12 source groups and shows AI minus human. References per claim are lower for AI in every group; zero-uptake is higher in every group; local family share is lower in every group. Point sizes encode paired sample count.

has 5.68×10^{534} nodes—a 535-digit count—although shared expression structure lets Lean store and check it without materializing that tree. Two other AI `norm_num` certificates have 281- and 427-digit expanded sizes; the largest human count has 11 digits. These numbers are not a cognitive score. They show a real compile-first route in which a tiny source episode denotes an astronomical repeated computation, while the human route identifies a reusable law. In complexity language, sharing compresses a formula into a circuit, but the human step does something different again: it changes the algorithmic description from instance expansion to an inductive invariant. The pair was selected for clarity; the aggregate family and uptake contrasts were fixed independently, and reverse examples exist.

3.4 Claim boundaries carry a short information pulse

Human conversational production regulates information in time: pauses and backchannels appear near changes in information rate (Bergey and DeDeo, 2024). If a named claim is an explanatory boundary rather than only syntax, it might likewise concentrate material that a predictive model finds locally unexpected.

The exploratory result is compatible with that prediction (figure 3). Across the 169 pairs in which both proofs contain eligible boundaries, the median eight-token boundary excess—post-boundary minus pre-boundary NLL, corrected by same-document control positions—is 0.334 nats per token human and 0.020 AI; the mean AI-minus-human source-cluster interval is $-.557$ to $-.276$. Removing the complete `have` header leaves medians 0.175 versus 0.036 (mean-difference interval $-.401$ to $-.130$). The effect is present in 4-, 8-, and more weakly 16-token windows, and disappears by 32 tokens. Thus it is a local pulse, not sustained global difficulty.

The model assigns AI documents much lower NLL overall (median 0.73 versus 1.15), a warning about style affinity or data overlap. A leave-one-source-out character bigram preserves both mean directions under source-cluster resampling, but its rank-based content contrast is marginal ($p = .093$). We infer neither creativity nor cognitive surprise. Nor does the pulse specifically validate the family classifier: within documents containing both kinds of claim, family and instance content pulses do not differ detectably (human $p = .62$; AI $p = .10$). At most, human boundaries mark a brief predictive discontinuity whose semantic interpretation remains model-dependent.

4 From provenance to objective

The horizon principle would be weak if it merely renamed the observed label. Its main content is that the difference should move when the objective moves. Several recent lines of work support this prediction.




Figure 3: Exploratory model-relative information around source `have` boundaries. Curves show negative log probability relative to the preceding 16 tokens; bands are proof-weighted 95% standard-error intervals. The model is Goedel-Prover-V2-8B in raw-code mode. The result may reflect training overlap, naming conventions, or prover-style affinity and is not a human surprisal measurement.

REFACTOR extracts heavily reused theorems from completed Metamath proofs (Zhou et al., 2024). Proof-Refactor and Lean Refactor explicitly optimize modularity, maintenance, checking cost, or version robustness after compilation (Fu et al., 2026; Lu et al., 2026). CircuitProver provides a closer intervention: accumulating verified lemmas and strategies across related Lean tasks reduces both proof length and checking time in ablation (Yang et al., 2026). Rtl2lean likewise asks a generator to build a layered reusable library (Lyu et al., 2026). At research-development scale, LeanMarathon makes an evolving formal and natural-language blueprint the shared state of a long-running autoformalization process; durable contracts and local repair support 258 proved declarations across its reported runs (Zhang et al., 2026b).

The ingredients have precedents: mined HOL Light inferences, LEGO-Prover, and DreamProver all turn intermediate results into later proving capacity (Kaliszyk and Urban, 2014; Wang et al., 2023; Zhang et al., 2026a). They also expose the retrieval constraint: persistent memory helps only if a later agent can select it; LeanSearch v2 reports a large gain from global premise retrieval (Gao et al., 2026). Our contribution is a common account of compile-first artifacts, refactored proofs, readable boundaries, and growing libraries, plus an objective that charges construction against future savings.

These systems do not prove that their artifacts match mature human libraries, and their metrics are not directly comparable with ours. They establish the more fundamental point: machines can be made to externalize persistent mathematical memory. The observed human–AI difference is therefore best read as a difference between present production regimes. The causal experiment now suggested is straightforward: assign the same base model the same theorem stream under (i) per-theorem compile reward and (ii) downstream cost, maintenance, and human-retrieval reward. Measure dead binders, claim reach, API survival, later citation, repair cost, and reader comprehension. If the horizons do not separate, our account fails.

The distinction also clarifies DeDeo’s discussion of “alien lemmas.” His argument draws on proof-complexity and average-case results to contrast hard-to-prove truths with a mathematical practice that seems to select truths with good reasons (DeDeo, 2024b). An AI may discover a lemma that dramatically compresses later certificates yet remains unrecognizable to human mathematical ontology. Such a lemma has high machine amortized value and low human interface value. It is neither simply good nor bad. It exposes two horizons with different consumers. The correspondence problem then becomes an engineering and epistemic demand: construct

a bridge that preserves the machine compression while letting human readers track why the abstraction matters (DeDeo, 2024a; DeDeo and Duede, 2026).

5 Cognitive implications: abstraction as memory allocation

The source-level difference suggests a cognitive interpretation more specific than “humans are abstract” and more defensible than reading mental states directly from code. Constructing a proof requires allocating limited representational capacity. A solver can preserve the past by keeping facts in the local context, by naming a relation that summarizes many facts, or by retrieving a library theorem that packages an older derivation. These are memory operations at different scales. The horizon determines which compression is worth performing.

At the scale of one search episode, an intermediate claim can function like a checkpoint. It freezes a successful subcomputation, constrains the remaining goal, and provides the next tactic with an enlarged context. Generic labels are adequate because identity can be recovered from position. Dead elaborated binders are not necessarily mistakes on this account. They are traces of a search process whose utility was exhausted before the final certificate. Removing them may improve the artifact while making the original generation harder to understand, just as deleting a scratch calculation produces a cleaner solution but erases evidence about discovery.

At a longer scale, an abstraction is a lossy compression chosen for transfer. The statement forgets derivational detail while preserving a relation expected to matter again. A descriptive name supplies a retrieval cue; a proof boundary asserts that the interior can be temporarily ignored; a general signature enlarges the class of future contexts in which the package applies. This compression is not free. Choosing the wrong invariant can increase working-memory demand, and exporting too many lemmas makes the library harder to search. Long-horizon construction is therefore not maximization of lemma count. It is selection of a small basis whose elements remain useful after the local context has vanished.

The aggregate near-equality of multi-use term binders is theoretically important here. It hides opposite channel components: AI binders disappear more often, but survivors are more often multiply represented. Kernel sharing and communicative abstraction differ. Automation can expand a generic hypothesis into many term occurrences without ever identifying a concept for a reader. Conversely, a well-chosen explanatory lemma can deserve a name even when invoked once in the current theorem, because its intended reuse lies in another proof or another mind. Any cognitive account based only on the certificate DAG will confuse these cases. Source references expose deliberate addressability; the term records compiled dependencies. Human–AI divergence can reverse between these planes.

The local surprisal pulse offers a tentative temporal signature of this selection. A useful boundary often arrives when prior manipulations have made a new description possible: monotonicity, invariance, symmetry, or a change of representation. The words immediately after the boundary can therefore be hard to predict from nearby syntax even though the resulting claim makes the rest of the argument easier. This connects to an account on which explanations of necessary truths mark search-simplifying steps or sites of cognitive load (Kardeş and DeDeo, 2025). In conversational terms, the constructor briefly spends information to renegotiate common ground. The pulse’s disappearance after 16 tokens is consistent with a boundary event rather than generally ornate prose. But this interpretation must remain provisional until human reading time, recall, and transfer are measured directly.

This distinction echoes geometry-expertise models in which abstract key steps and perceptual chunks guide solving rather than stepwise planning (Koedinger and Anderson, 1990). A source proof can be a search trace, a pedagogical artifact, or both. Humans often refactor scratch work into a proof idea; present generators are commonly evaluated before that second pass. A cognitively serious comparison should align stages—first attempts with first attempts, refactored artifacts with refactored artifacts—and observe how representations change between them. The horizon principle predicts that the largest human–AI gap will occur not in local inference, where both can exploit formal context, but in deciding what part of a successful episode should become shared memory.

6 Library-amortized proof complexity

Ordinary proof complexity asks for the size of a certificate for one formula in a fixed proof system (Cook and Reckhow, 1979; Krajíček, 1995). This leaves library construction outside the objective. Let $\Phi = (\varphi_1, \dots, \varphi_T)$ be a stream of targets. Let L be an ordered, acyclic sequence of auxiliary declarations: each signature and body

must check in the base environment plus earlier members of L , and $|L|$ charges both. If P_t checks φ_t in that environment plus earlier checked target declarations and $|P_t \mid L, P_{<t}|$ is its residual certificate size, define

$$\text{LAPC}_\lambda(\Phi) = \min_{L, P_1, \dots, P_T} \left[\lambda|L| + \sum_{t=1}^T |P_t \mid L, P_{<t}| \right], \quad (2)$$

with $\lambda \geq 1$. Thus neither a circular declaration nor an unpriced advice oracle is allowed. Unlike nonuniform proof systems with advice (Beyersdorff et al., 2011), L is stream-specific, checked in the same environment, and fully charged. The factor λ prices the design and maintenance of shared structure. A richer version can price names, dependency fragility, retrieval, and human validation separately. An online version charges when each declaration is introduced and compares cumulative cost with the best hindsight library.

Let $C(\varphi)$ be the shortest proof of one target without a new library. Because L may be empty, $\text{LAPC}_\lambda(\Phi) \leq \sum_t C(\varphi_t)$. The gap

$$G_\lambda(\Phi) = \sum_t C(\varphi_t) - \text{LAPC}_\lambda(\Phi) \quad (3)$$

is the compression attributable to shared mathematical infrastructure at the chosen interface price. At $\lambda = 1$ the objective is a restricted two-part description length: library plus residual certificates. It is not Kolmogorov complexity; the language is fixed to checked Lean artifacts and retrieval cost can be charged explicitly.

Even a syntactic shadow of this objective is computationally hard. If certificates are strings and library entries are restricted to acyclic macros, minimizing library plus residual size becomes the smallest-grammar problem, which is NP-hard and hard to approximate closely (Charikar et al., 2005). Typed declarations add validity, elaboration, retrieval, and semantic equivalence to that selection problem. Thus a long-horizon agent must solve not only proof search but a difficult representation-design problem about which repeated structure deserves a boundary.

The break-even law is immediate but useful. Ignoring interactions, a declaration of charged size b that saves s symbols in each of k residual proofs pays exactly when $ks > \lambda b$; with temporal discounting, ks becomes $\sum_i \gamma^{d_i} s_i$. The amortization horizon is therefore not a metaphor but the point at which checked shared structure beats repeated local construction.

Equation (2) changes what counts as an optimal proof. A declaration that lengthens the first episode may lower total cost if it compresses later proofs. A generic helper that is short but impossible to retrieve may fail to pay back its interface cost. A benchmark of independent theorems with a reset context implicitly sets future reuse to zero and selects a compile-first policy. Even pass@ k with enormous search budgets remains a measure of finding current certificates, not of learning a mathematical language in which later certificates become cheap.

The connection to circuit sharing and extension variables is instructive but limited. Naming a repeated subexpression turns tree accounting into DAG accounting; introducing a reusable theorem resembles adding a subroutine or grammar rule. Yet Lean claims are dependent types, tactics alter elaboration, definitions compute, and public theorem statements carry semantic and maintenance costs. We do not claim an equivalence with Extended Frege. The robust lesson is accounting level: single-certificate complexity cannot see investments whose return lies across certificates.

This produces new empirical questions at the boundary of complexity and cognition. How large can online library regret be for a bounded-context agent? Which theorem orderings permit discovery of the hindsight-optimal abstractions? Does a human-readable interface impose only polynomial overhead on a machine-efficient lemma basis, or can correspondence itself be exponentially costly? DeDeo’s “hard proofs” dilemma suggests that mathematical culture may be understood as a selection process over theorem streams and representations, concentrating attention where useful proofs and reasons coincide (DeDeo, 2024b). AI changes both the search distribution and the consumers for whom that coincidence must hold.

7 What the result does and does not establish

The result is structural, not causal. Here AI emits more local state, but less is adopted, more vanishes under elaboration, and names are more generic; conditional duration does not robustly differ, while surviving binders can fan out. Global reuse is unobserved. Provenance is observational; the census misses other constructs and implicit use; binder occurrence is representation-specific; and names, spans, and surprisal are proxies. Reader experiments remain necessary (Kambhamettu et al., 2026); the topology null and scope correction are in the appendix.

References

- AI-MO. NuminaMath-LEAN: A large-scale dataset of competition problems formalized in Lean 4. Hugging Face dataset, <https://huggingface.co/datasets/AI-MO/NuminaMath-LEAN>, 2025.
- Claire Augusta Bergey and Simon DeDeo. From “um” to “yeah”: Producing, predicting, and regulating information flow in human conversation, 2024. arXiv:2403.08890.
- Olaf Beyersdorff, Johannes Köbler, and Sebastian Müller. Proof systems that take advice. *Information and Computation*, 209(3):320–332, 2011. doi: 10.1016/j.ic.2010.11.006.
- Moses Charikar, Eric Lehman, Ding Liu, Rina Panigrahy, Manoj Prabhakaran, Amit Sahai, and Abhi Shelat. The smallest grammar problem. *IEEE Transactions on Information Theory*, 51(7):2554–2576, 2005. doi: 10.1109/TIT.2005.850116.
- Katherine M. Collins, Simon Frieder, Jonas Bayer, Jacob Loader, Jeck Lim, Peiyang Song, Fabian Zaiser, Lexin Zhou, Shanda Li, Sam Looi, Joshua B. Tenenbaum, Umang Bhatt, Adrian Weller, Jose Hernandez-Orallo, Cameron E. Freer, Valerie Chen, and Ilia Sucholutsky. Characterizing initial human–ai proof formalization workflows, 2026. arXiv:2606.04273.
- Stephen A. Cook and Robert A. Reckhow. The relative efficiency of propositional proof systems. *Journal of Symbolic Logic*, 44(1):36–50, 1979. doi: 10.2307/2273702.
- Simon DeDeo. Alephzero and mathematical experience. *Bulletin of the American Mathematical Society*, 61(3): 375–386, 2024a. doi: 10.1090/bull/1824.
- Simon DeDeo. Hard proofs and good reasons, 2024b. arXiv:2410.18994.
- Simon DeDeo and Eamon Duede. A correspondence problem for mathematical proof, 2026. arXiv:2603.13680.
- Kosta Došen. Identity of proofs based on normalization and generality, 2004. arXiv:math/0208094.
- Catarina Dutilh Novaes. A dialogical conception of explanation in mathematical proofs. In Paul Ernest, editor, *The Philosophy of Mathematics Education Today*, pages 81–98. Springer, 2018. doi: 10.1007/978-3-319-77760-3_5.
- Yiming Fu, Peixuan Liu, Zichen Wang, and Kun Yuan. Proof-refactor: Refactoring generated formal proofs into modular artifacts, 2026. arXiv:2606.03743.
- Guoxiong Gao, Zeming Sun, Jiedong Jiang, Yutong Wang, Jingda Xu, Peihao Wu, Bryan Dai, and Bin Dong. LeanSearch v2: Global premise retrieval for Lean 4 theorem proving, 2026. arXiv:2605.13137.
- Cezary Kaliszyk and Josef Urban. Learning-assisted theorem proving with millions of lemmas. In *International Symposium on Symbolic and Numeric Algorithms for Scientific Computing*, pages 378–384, 2014. arXiv:1402.3578.
- Hita Kambhamettu, Will Crichton, Sean Welleck, Harrison Goldstein, and Andrew Head. Making written theorems explorable by grounding them in formal representations, 2026. arXiv:2604.02598.
- Gülce Kardeş and Simon DeDeo. Explaining necessary truths. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 47, 2025. arXiv:2502.11251.
- David Kirsh and Paul Maglio. On distinguishing epistemic from pragmatic action. *Cognitive Science*, 18(4): 513–549, 1994. doi: 10.1207/s15516709cog1804_1.
- Kenneth R. Koedinger and John R. Anderson. Abstract planning and perceptual chunks: Elements of expertise in geometry. *Cognitive Science*, 14(4):511–550, 1990. doi: 10.1207/s15516709cog1404_2.
- Jan Krajíček. *Bounded Arithmetic, Propositional Logic and Complexity Theory*. Cambridge University Press, 1995. doi: 10.1017/CBO9780511529948.

- Jill H. Larkin and Herbert A. Simon. Why a diagram is (sometimes) worth ten thousand words. *Cognitive Science*, 11(1):65–100, 1987. doi: 10.1111/j.1551-6708.1987.tb00863.x.
- Lean Developers. *The Lean Language Reference: The Type System and Propositions*, 2026. URL <https://lean-lang.org/doc/reference/latest/The-Type-System/>.
- Falk Lieder and Thomas L. Griffiths. Resource-rational analysis: Understanding human cognition as the optimal use of limited computational resources. *Behavioral and Brain Sciences*, 43:e1, 2020. doi: 10.1017/S0140525X1900061X.
- Yong Lin, Shange Tang, Bohan Lyu, Ziran Yang, Jui-Hui Chung, Haoyu Zhao, Lai Jiang, Yihan Geng, Jiawei Ge, Jingruo Sun, Jiayun Wu, Jiri Gesi, Ximing Lu, David Acuna, Kaiyu Yang, Hongzhou Lin, Yejin Choi, Danqi Chen, Sanjeev Arora, and Chi Jin. Goedel-prover-v2: Scaling formal theorem proving with scaffolded data synthesis and self-correction, 2025. arXiv:2508.03613.
- Øystein Linnebo. Generality explained. *The Journal of Philosophy*, 119(7):349–379, 2022. doi: 10.5840/jphil2022119725.
- Jialin Lu, Soonho Kong, Rodrigo Stehling, Kaiyu Yang, Zhangyang Wang, Weiran Sun, and Wuyang Chen. Lean refactor: Multi-objective controllable proof optimization via agentic strategy search, 2026. arXiv:2605.20244.
- Hongqin Lyu, Junxing Dong, Yonghao Wang, Zhiteng Chao, Tiancheng Wang, and Huawei Li. Rtl2lean: Automated RTL-to-Lean translation with hierarchical theorem generation and lemma reuse, 2026. arXiv:2607.16855.
- Numina and Kimi Teams. Kimina-Prover Preview: Towards large formal reasoning models with reinforcement learning, 2025. arXiv:2504.11354.
- Yehuda Rav. Why do we prove theorems? *Philosophia Mathematica*, 7(1):5–41, 1999.
- Claude E. Shannon. Coding theorems for a discrete source with a fidelity criterion. In *IRE National Convention Record*, volume 4, pages 142–163, 1959.
- Scott Viteri and Simon DeDeo. Epistemic phase transitions in mathematical proofs. *Cognition*, 225:105120, 2022. doi: 10.1016/j.cognition.2022.105120.
- Haiming Wang, Huajian Xin, Chuanyang Zheng, Lin Li, Zhengying Liu, Qingxing Cao, Yinya Huang, Jing Xiong, Han Shi, Enze Xie, Jian Yin, Zhenguo Li, Heng Liao, and Xiaodan Liang. LEGO-Prover: Neural theorem proving with growing libraries, 2023. arXiv:2310.00656.
- Zachary Wojtowicz and Simon DeDeo. Undermining mental proof: How AI can make cooperation harder by making thinking easier. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(2):1592–1600, 2025. doi: 10.1609/aaai.v39i2.32151.
- Ziyi Yang, Wenji Fang, Chen Chen, Zhiyao Xie, and Hongce Zhang. CircuitProver: Agentic Lean 4 theorem proving with reusable circuit proof library for hardware verification, 2026. arXiv:2607.27259.
- Jiajie Zhang and Donald A. Norman. Representations in distributed cognitive tasks. *Cognitive Science*, 18(1): 87–122, 1994. doi: 10.1207/s15516709cog1801_3.
- Youyuan Zhang, Jialiang Sun, Hangrui Bi, Chuqin Geng, Wenjie Ma, Zhaoyu Li, and Xujie Si. DreamProver: Evolving transferable lemma libraries via a wake-sleep theorem-proving agent, 2026a. arXiv:2604.26311.
- Yuanhe Zhang, Yuekai Sun, Taiji Suzuki, Jason D. Lee, and Fanghui Liu. LeanMarathon: Toward reliable AI co-mathematicians through long-horizon Lean autoformalization, 2026b. arXiv:2606.05400.
- Jin Peng Zhou, Yuhuai Wu, Qiyang Li, and Roger Grosse. REFACTOR: Learning to extract theorems from proofs, 2024. arXiv:2402.17032.

A Reproducibility and robustness

All reported paired source results are produced by `code/paired_horizon.py`; scope-aware term use by `code/extract_binder_use.py` and `code/analyze_binder_use.py`; and surprisal by `code/prepare_surprisal.py`, `code/token_surprisal.cpp`, `code/ngram_surprisal.py`, and `code/analyze_surprisal.py`. Machine-readable summaries live in `results/horizon/`. Seeds and inclusion rules are serialized with the estimates.

The source census uses `code/ExtractHaveRanges.lean` to parse both the complete tactic that constructs each claim and the nearest enclosing tactic-sequence scope before searching for later references. This prevents an outer same-name hypothesis used in `have h : P := f h` from being credited to the newly introduced `h`, and prevents a same-spelled token in a sibling branch from being counted after the binder has gone out of scope. Scope clipping removes 1,068 reference-like tokens from 442 human claims and 1,157 from 568 AI claims; the older unscoped rates were 1.43/0.87 uses per claim, compared with the corrected 1.35/0.82. The current parser aligns 39,646/40,626 candidates. Restricting to the 3,511 theorem pairs with successful parsing and complete scanner/parser agreement gives human/AI uses per claim 1.35/0.82, zero-uptake 22.2%/36.9%, and multi-uptake 24.0%/10.3%; the four source-cluster intervals for the between-track differences exclude zero.

The supply/selectivity decomposition uses proof-body tokens as a second denominator. AI-minus-human source-cluster intervals, in claims per 100 tokens, are .304-.605 for adoption at least once, -.179--.132 for two-or-more references, -.514--.368 for non-placeholder naming, and -.073--.036 for generalized claims. The elementary boundary-padding construction explains why raw count cannot identify clarity: before any term $p : P$, insert distinct unused declarations `have junk_i : True := by trivial`. This produces P for every finite number of insertions while raising source boundary count arbitrarily; fate and construction cost, unlike count, reveal the padding. At minimum reference counts three and four, human/AI supply is .221/.109 and .124/.052; all 12 source groups favor humans at thresholds two through four. Among the 834 length-matched pairs, the AI advantage remains at zero references (interval .372-.694), is unresolved at one (-.004-.311), and reverses at two (-.151--.004) through four. Character normalization preserves every reference-count direction.

For a stricter episode proxy, each name token is assigned to the innermost later named-claim construction containing it; all remaining in-scope tokens form one conservative residual site. Human claims reach multiple such sites more often (9.13%/6.26%; interval -.049--.018). Absolute supply is .179/.194 per 100 whitespace tokens (-.004-.030), also unresolved in the length match, but .375/.440 per 1,000 characters favors AI (.023-.093). We therefore make no absolute multi-site provenance claim. The site metric undercounts unnamed tactic episodes, but prevents two mentions inside one construction from masquerading as two consumers.

Those sites also define a source construction DAG: nodes are parser-matched named claims, and an edge connects an earlier claim to a later claim whose complete construction references it. Pooled edges per node are .325 human and .296 AI, named-claim branchpoint shares .052/.042, and longest-chain length per node .115/.102; all difference intervals narrowly exclude zero. More decisively, in the 206 equal-positive-node pairs, edge rates remain .380/.339 (-.092--.010) and branchpoint shares .088/.075 (-.030--.003), while maximum chain length is unresolved. AI proofs contain more nodes and, in absolute terms, more edges per proof, but the node-normalized graph is sparser and does not show greater seriality at equal node count. Restricting to names declared once per proof strengthens the edge (.386/.350) and branchpoint (.067/.055) directions; both intervals exclude zero. Equal node count makes later-node opportunities identical within each pair, and both matched directions survive every leave-one-source deletion; the branching contrast is not an earlier-position or one-source effect.

Source intervals resample 12 source groups with replacement and recompute pooled claim ratios. Every source group has lower explicit uses per AI claim, higher AI zero-uptake, and lower AI long-reach share. The production term alignment’s unique-name subset contains 9,226 human and 14,706 AI claims; zero-use shares are 9.5% and 21.8%, respectively. The position-matched family-retention result also survives this unique-name restriction (202 human and 167 AI proof sides); the human within-track multi-use association remains positive, while the AI association and the direct multi-use provenance interaction remain unresolved. The surprisal boundary contrast is present for four-, eight-, and more weakly 16-token windows, and absent at 32. Excluding the complete claim header preserves the four- and eight-token effects. In the leave-one-source-out bigram, which excludes every document in the held-out proof’s coarse source group from training, the raw eight-token rank difference remains ($p = .020$) but content-only rank evidence weakens ($p = .093$); their mean source-cluster intervals are -.426 to -.081 and -.540 to -.062.

The adoption-duration decomposition guards against a changing temporal ruler. For every adopted claim, token reach is the number of lexical tokens from the end of its complete construction tactic to its last explicit reference. Opportunity-normalized reach divides crossed later claim boundaries by the number available before the scope or same-name stopping point and excludes declarations with no later boundary. Human token reach is higher in all 12 source groups. In contrast, the conditional chance of crossing at least one available boundary is 83.6%/85.0% (AI-minus-human source-cluster interval -.005 to .028), and normalized reach is 63.7%/62.0% (interval -.046 to .003). In the length-matched subset, both conditional-duration intervals cross zero; under equal positive claim counts, normalized duration is 64.3%/64.2%. The claim-density ruler alone would conceal both null results.

The local functional check on the family classifier holds separately for its two equivalent spellings. Within eligible

proofs, the multi-use uplift for a pre-type-binder family is 20.7 points human and 29.1 points AI; for a leading-universal family it is 30.8 and 38.6 points. All four source-cluster intervals exclude zero. Thus the result is not created by unioning a human-favored and an AI-favored surface form.

Claim position is a second mechanical alternative: an early family simply has more remaining text in which to be referenced. Within each proof we therefore use minimum-cost one-to-one matching on claim index normalized by the proof’s final claim index. With no caliper this yields 577 human and 409 AI claim pairs; with an absolute position-gap caliper of .25, it yields 394 and 348 pairs in 248 and 199 proofs. In the calipered comparison, family-minus-instance adoption is 15.0 points human and 20.6 AI, while multiple-reference uplift is 26.8 and 37.0 points. All intervals exclude zero. The explicit-reference count itself is less stable for humans, but the two binary functional outcomes are not explained by simple exposure time.

We also attempted a stronger test of the naming interpretation. In 229 pairs with two to eight unique, explicitly typed claims per proof, each naturalized name had to retrieve its own type from the other types in that proof using Goedel-Prover mean embeddings. Mean percentile rank is .528 human and .519 AI (AI-minus-human source-cluster interval $-.020$ to $.006$); chance-adjusted top-one accuracy is likewise unresolved. Last-token pooling gives the same null direction, and among 61 equal-claim-count pairs mean-pooling top-one accuracy is 40.7% human and 39.9% AI. Because this causal model was not trained as a contrastive encoder, even a positive result would have been exploratory. The null is sufficient for the present inference: vocabulary size and generic-form frequency do not establish that human names are semantically better retrieval keys.

The preselected 312-pair current-version sample closely preserves the source contrasts. Human/AI claim densities are 1.94/3.16 per 100 tokens in the sample versus 1.96/3.10 in the full corpus; explicit uses per claim are 1.33/0.80 versus 1.35/0.82; and zero-uptake shares are 20.8%/35.8% versus 22.2%/36.9%. All 12 source groups remain represented. Current-version attrition reduces this to 298 complete binder-use pairs but changes the profile little.

Paired does not mean independent: 181 structured proof values are token-identical (180 also have identical source bodies), and 230 have whitespace-token similarity at least .90 after comments and strings are removed. Excluding all 230 leaves 3,400 pairs and preserves the per-claim uptake directions and lexical supply crossover. Among the 172 retained equal-node pairs, human edge and branchpoint advantages survive; chain depth remains unresolved.

The source direction also survives a coarse automation sensitivity. For each of `linarith`, `nlinarith`, `norm_num`, `omega`, `ring`, and `simp`, we separately retain pairs in which both sides invoke the tactic and pairs in which neither does. Across all 12 strata, AI claims have fewer explicit references per claim and a higher zero-uptake share. This does not equate full tactic workflows, but rules out any one of these common families as a sufficient explanation.

B Exact counting without tree expansion

The production extractor never materializes the expanded proof tree. For each expression node e , it computes a summary $S(e) = (N_e, O_e, B_e)$: the number N_e of nodes in the fully expanded syntax tree, a finite map O_e from free de Bruijn depth to occurrence count, and the ordered local binders decoded below e . Application and constructor nodes add child summaries once per child occurrence. Entering a binder records $O_e(0)$ as its use count, drops that coordinate, and shifts the remaining coordinates down by one. Because N_e is an arbitrary-precision natural, shared subexpressions can contribute exponentially without expansion. An explicit work stack makes both the semantic summary and raw-tree recurrence stack safe.

Memoization requires one subtle restriction. Lean expression hashes are insensitive to some alpha-renaming, while this audit must recover the source-facing binder names. Caching a binder-bearing summary can therefore transplant a hygienic name between alpha-equivalent scopes. The production traversal caches only summaries with no decoded binders; de Bruijn occurrence maps remain sufficient for scope-correct composition. Every extraction records the template, toolchain, manifest, and corpus hashes.

We checked equivalence rather than trusting the optimization. A prespecified audit unions the first 50 lexicographic legacy-success pairs, the 50 largest legacy-success roots, and ten seeded random pairs from each of 12 source groups. This yields 215 distinct pairs (430 tasks), with a largest legacy root of 97,350,243 nodes. All 430 memoized records equal the recursive legacy records after replacing only the temporary-work-tag substring inside Lean-generated hygienic names; 427 already agree byte-for-byte as parsed JSON. The same canonical comparison over all 7,186 tasks completed by the legacy traversal is exact, and the memoized traversal additionally recovers all 74 legacy timeouts. The audit selection, mismatches, hashes, and summary are serialized in `results/horizon/binder_memo_equivalence.json` and `results/horizon/binder_memo_full_equivalence.json`.

The recurrence example in the main text is the largest representational extreme in this complete census. Fully expanded tree counts reach 11 decimal digits on the human side and 535 on the AI side; four AI tasks exceed 20 digits and three exceed 100, versus none for humans. These counts are useful precisely because they expose the formula represented by a shared expression DAG. They are not heap size, checking time, semantic proof complexity, or evidence that Lean constructs the tree.

C Executable representation regressions

Three small Lean regressions make the invariance corrections independently checkable. First, `code/test_legacy_binder.py` constructs the same local family once with pre-type binders and once with an explicit universal type. Under Lean 4.15 both elaborate to a named `letFun` binder used once. The source classifier must place both in its generalized class.

Second, `code/test_binder_toolchain_shift.py` elaborates the same real corpus theorem under the production and current snapshots. Its source claim `h1` is a decoded `have_fun` in Lean 4.15 and a nondependent `let` in Lean 4.33, with one occurrence in both. This is why raw node kind cannot define retention across versions.

Third, `code/test_certificate_representation.py` builds a well-typed `Nat` expression $e_{n+1} = \text{let } x := e_n; x + x$, invokes Lean’s zeta reducer, and asks Lean to infer both types. At depth 14 the raw and reduced expressions both have type `Nat`, their types are definitionally equal, and their measured sizes are 99 and 65,533 nodes. These are regressions, not benchmarks: their purpose is to prevent later analyses from silently treating a representation choice as a fact about reasoning.

D Why the old network result is not the new result

The prior project replicated Viteri and DeDeo’s fixed-tail power-law results and modeled epistemic phase transitions across Coq and Lean. Those are properties of declared network constructions. They do not supply a representation-invariant authorship signature. Tail exponents move when the minimum degree or expansion boundary moves; human–AI directions can reverse under proof-value normalization; and modeled theorem belief is close to saturation in both tracks.

The scope audit strengthens the reason for restraint. A de Bruijn index identifies a binder by distance in its local expression, not by a global name. Interning equal raw indices across scopes turns distinct variables into one graph node. The corrected traversal carries a scope environment and replaces each entered binder with a fresh free-variable identity before hashing expressions. The regression proof has two independent identity lambdas; the old representation produces one shared variable leaf, while the corrected representation produces two. None of the main analyses in this paper uses a raw-term reuse tail.

The scope-correct rerun elaborates 302 of 312 tasks on each side; applying the independent binder-task status filter leaves 298 complete pairs. AI terms are slightly larger (median paired node difference 88.5; source-cluster interval 0–243) and have a slightly larger visit-to-unique-node ratio (difference 0.0068; 0.0023–0.0112). But the proposed authorship signal in high-outdegree reuse remains absent: maximum outdegree and the fixed-10 tail exponent have intervals crossing zero. These graph quantities remain representation-specific diagnostics, not measures of cognitive abstraction.

E Rate–distortion form of the horizon principle

Let an episode $E \sim p(e)$ include the derivation history available at consolidation time, let $U \sim q_\gamma$ be a future use drawn from a horizon-weighted task distribution, and let an interface encoder choose $Z \sim p(z | E)$. For consumer c , let $d_c(U; E, Z) \geq 0$ charge downstream search, repair, retrieval failure, or loss of the ability to track the proof idea, and define horizon-weighted distortion

$$D_c(E, Z) = \mathbb{E}_{U \sim q_\gamma} [d_c(U; E, Z)].$$

A consumer-specific interface frontier is

$$R_c(\delta) = \min_{p(z|E): \mathbb{E} D_c(E, Z) \leq \delta} I(E; Z).$$

For a deterministic encoder, the rate is the information retained in the chosen interface. A shared encoder is assumed: $I(E; Z)$ does not itself charge for inventing or communicating the codebook, a cost restored by the explicit $|L|$ term in [equation \(2\)](#). A short horizon changes q_γ toward the immediate goal, making positional scratch state cheap and distant transfer nearly irrelevant. A different consumer changes D_c : an opaque lemma may have low distortion for a prover with effective retrieval and high distortion for a reader who cannot connect it to the proof idea.

Three claims follow. Consumer-invariance is the zero-distortion case: a functional statistic must identify artifacts in the same response class. Generality enlarges the support over which a code can pay, while naming can lower retrieval loss. And [equation \(2\)](#) is a checked finite-stream analogue: library size is rate, residual certificate cost is downstream distortion. The human–AI question is not who compresses, but which future-use distribution and distortion govern it.

A proposed causal benchmark. A causal test samples theorem sequences, randomizing agents between compile-first and persistent-library objectives with fixed library, target order, Lean version, and budget. Charge current cost, reward later savings, repair, and blinded retrieval, and freeze declarations before future targets. Measure cumulative cost, hindsight regret, API survival, citations, dead binders, and delayed comprehension. A human/AI $\times$ objective crossover separates horizon from provenance.